

Explore Twitter Data

Lyn Nguyen

2022-12-07

Contents

ELT	1
Get <code>text</code> and <code>tweet_id</code> only.	3
Clean up <code>users</code>	3
EDA	5
experiment_group / in_reply_to_screen_name	5
Tweet	6
text	6
Tweet Popularity	8
retweets & experiment_group	8
tweet_likes & experiment_group	9
User	10
screen_name	10
created_at	11
description	13
location	13
ymd & dow	13
User Popularity	17
favourites_count	17
followers_count	17
verified	17
EDA - with annotations	18

ELT

This script uses output from `analysis-of-public-opinion/scrapper.py`. Ultimately, we keep data pulled on Dec

```

# created_at to date and day of week
test = head(tweets1)
dow <- substr(test$created_at, 1, 3)
month_day <- substr(test$created_at, 5, 10)
time<- substr(test$created_at, 12, 19)
yr <- substr(test$created_at, 26, 30)

ymd <- as.Date(paste0(month_day, yr), format = "%b %d %h:%m:%s %Y")

# as.Date(test$created_at, format = "%a %b %d %h:%m:%s +0000 %Y")
tweets1 <- tweets1 %>% mutate(dow = substr(created_at, 1, 3)
                             , month_day = substr(created_at, 5, 10)
                             , time = substr(created_at, 12, 19)
                             , yr = substr(created_at, 26, 30),
                             , ymd = as.Date(paste0(month_day, yr), format = "%b %d %Y"))
tweets2 <- tweets2 %>% mutate(dow = substr(created_at, 1, 3)
                             , month_day = substr(created_at, 5, 10)
                             , time = substr(created_at, 12, 19)
                             , yr = substr(created_at, 26, 30),
                             , ymd = as.Date(paste0(month_day, yr), format = "%b %d %Y"))
tweets3 <- tweets3 %>% mutate(dow = substr(created_at, 1, 3)
                             , month_day = substr(created_at, 5, 10)
                             , time = substr(created_at, 12, 19)
                             , yr = substr(created_at, 26, 30),
                             , ymd = as.Date(paste0(month_day, yr), format = "%b %d %Y"))
                             , tweet_id_char = as.character(as.numeric(tweet_id)))
tweets4 <- tweets4 %>% mutate(dow = substr(created_at, 1, 3)
                             , month_day = substr(created_at, 5, 10)
                             , time = substr(created_at, 12, 19)
                             , yr = substr(created_at, 26, 30),
                             , ymd = as.Date(paste0(month_day, yr), format = "%b %d %Y"))
                             , tweet_id_char = as.character(as.numeric(tweet_id)))

summary(tweets1$ymd)

```

```

##           Min.         1st Qu.         Median         Mean         3rd Qu.         Max.
## "2022-11-26" "2022-12-01" "2022-12-01" "2022-12-01" "2022-12-03" "2022-12-03"

```

```
summary(tweets2$ymd)
```

```

##           Min.         1st Qu.         Median         Mean         3rd Qu.         Max.
## "2022-11-26" "2022-12-01" "2022-12-01" "2022-12-01" "2022-12-03" "2022-12-03"

```

```
summary(tweets3$ymd)
```

```

##           Min.         1st Qu.         Median         Mean         3rd Qu.         Max.
## "2022-11-27" "2022-12-01" "2022-12-02" "2022-12-02" "2022-12-03" "2022-12-05"

```

```
summary(tweets4$ymd)
```

```

##           Min.         1st Qu.         Median         Mean         3rd Qu.         Max.
## "2022-11-28" "2022-12-01" "2022-12-01" "2022-12-01" "2022-12-03" "2022-12-03"

```

tweets1.csv has data from 11/26/2022 but only cnn as liberal source. tweet2.csv: 11/26- 12/3 but only cnn as liberal source tweet3.csv: 11/28- 12/3 liberal sources has cnn, npr, msnbc, nytimes, tweet4.csv: 11/28- 12/3 but only cnn as liberal source

tweets1 and tweets2 have 814 fields total, but only 468 unique.

```
master <- rbind(tweets3, tweets4) %>% select(-experiment_id) %>% distinct()
```

master has 471 points, but length(unique(master\$tweet_id)) has 468 points. Where is the 3 difference? Since tweet4 hit the api after tweet3, some has updated values. For example tweet_id “1598304394931412992” has 0 like in tweet3 but 1 like in tweet 4. If there is duplicate in tweet_id, we will keep the one with the higher index.

```
master <- master %>% mutate(tweet_id_char = as.character(as.numeric(tweet_id)))
master_tweet_id <- master$tweet_id_char
dup_master <- master_tweet_id[duplicated(master_tweet_id) == T]

print("The duplicated tweet_ids are:")
```

```
## [1] "The duplicated tweet_ids are:"
```

```
dup_master
```

```
## [1] "1598304394931412992" "1598277959223083008" "1598411537667874816"
```

3 tweets are duplicated because they have updated “likes” count.

```
dup_val1 <- master[master$tweet_id == 1598304394931412992, ][2,]
dup_val2 <- master[master$tweet_id == 1598277959223083008, ][2,]
dup_val3 <- master[master$tweet_id == 1598411537667874816, ][2,]

m <- master %>% filter(!tweet_id %in% dup_master)
master <- rbind(m, dup_val1, dup_val2, dup_val3) %>% arrange(tweet_id) # in ascending tweet_id order

# write.csv(master, "prelim_data/tweets_master_dec5dec6.csv")
```

Get text and tweet_id only.

Madelaine will use this file in SageMaker. Need to keep row orders for annotation output.

```
tweet_text <- master %>% select("tweet_id", "text") %>% distinct() #468
# write.csv(tweet_text, "prelim_data/tweet_text_only.csv")
```

Clean up users.

```
# What user_id in user3 that's not in user4?
user4_id <- users4$user_id
user3_id <- users3$user_id
```

```

user3_id_only <- setdiff(user3_id, user4_id) #ids in user3 that's not in user4 - 118 total
user3_profiles_only <- users3 %>% filter(user_id %in% user3_id_only)

all_users <- rbind(user3_profiles_only, users4) %>% mutate(user_id_char = as.character(as.numeric(user_id)))

# merge users and tweets ==
# rename overlap column names
tweet_cnames <- colnames(master)
colnames(master) <- c("experiment_group", "text", "tweet_id", "tweet_likes", "retweets", "tweet_created_at",
                     "in_reply_to_screen_name", "screen_name", "dow", "month_day", "time", "yr", "ymd", "tweet_id")
author_cnames <- colnames(all_users)

final_tweets <- left_join(master %>% select(-c(screen_name)), all_users, by = "user_id")

write.csv(final_tweets, "data/master.csv")

```

There are 459 unique authors for these 468 tweets.

EDA

final_tweets have 25 columns and 468 observations (tweets).

```
glimpse(final_tweets)
```

```
## Rows: 468
## Columns: 25
## $ experiment_group      <chr> "msnbc", "msnbc", "msnbc", "msnbc", "msnbc", "~
## $ text                  <chr> "@MSNBC @MaddowBlog 'Simpleton's defense'? Yo~
## $ tweet_id              <dbl> 1.596988e+18, 1.596993e+18, 1.596997e+18, 1.59~
## $ tweet_likes           <int> 4, 0, 0, 2, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0~
## $ retweets              <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0~
## $ tweet_created_at      <chr> "Sun Nov 27 22:01:59 +0000 2022", "Sun Nov 27 ~
## $ user_id               <dbl> 1.518750e+18, 3.202809e+09, 1.409157e+08, 1.93~
## $ in_reply_to_status_id <dbl> 1.596987e+18, 1.596987e+18, 1.596987e+18, 1.59~
## $ in_reply_to_user_id   <int> 2836421, 2836421, 2836421, 2836421, 2836421, 2~
## $ in_reply_to_screen_name <chr> "MSNBC", "MSNBC", "MSNBC", "MSNBC", "MSNBC", "~
## $ dow                   <chr> "Sun", "Sun", "Sun", "Sun", "Mon", "Mon", "Mon~
## $ month_day             <chr> "Nov 27", "Nov 27", "Nov 27", "Nov 27", "Nov 2~
## $ time                  <chr> "22:01:59", "22:22:27", "22:39:00", "23:13:38"~
## $ yr                    <chr> " 2022", " 2022", " 2022", " 2022", " 2022", "~
## $ ymd                   <date> 2022-11-27, 2022-11-27, 2022-11-27, 2022-11-2~
## $ tweet_id_char         <chr> "1596987727953924096", "1596992880002084864", ~
## $ created_at            <chr> "Tue Apr 26 00:33:21 +0000 2022", "Sat Apr 25 ~
## $ description           <chr> "No name", "People following me are president ~
## $ location              <chr> "", "Massachusetts, USA", "Washington, DC", "w~
## $ followers_count       <int> 8, 874, 375, 537, 5, 130, 28, 200, 15, 18, 91,~
## $ screen_name            <chr> "BigTex1022", "michael_favreau", "AlxHamilt~
## $ statuses_count        <int> 2333, 30060, 33016, 60763, 1102, 1636, 1637, 1~
## $ favourites_count       <int> 1941, 16373, 1061, 19861, 320, 586, 1414, 5190~
## $ verified              <chr> "False", "False", "False", "False", "False", "~
## $ user_id_char          <chr> "1518749825092788224", "3202808548", "14091571~
```

experiment_group / in_reply_to_screen_name


What is the share of replies to the 5 news sources? How do ('msnbc', 'cnn', 'npr', 'nytimes') compare to 'cnn'? - FoxNews make up 87% of our data points. When it comes to the student loan forgiveness discussion, the Department of Education has the least engagement from Twitter users, at only 1%.

```
liberal <- c('msnbc', 'cnn', 'npr', 'nytimes')
conservative <- c('foxnews')
```

```
source_count <- as.data.frame(table(final_tweets$in_reply_to_screen_name)) %>% mutate(Proportion = round(
source_count
```

```
##      Var1 Freq Proportion
## 1 usedgov    6      0.01
## 2 nytimes    8      0.02
## 3 CNN       12      0.03
## 4 NPR        13      0.03
## 5 MSNBC     21      0.04
## 6 FoxNews  408      0.87
```

```
barplot(source_count$Freq)
```



Tweet

text

Is tweet length a distinguishable characteristic for the experiment groups? Within the liberal groups, most of NPR replies have over 40 words. NYTimes's reply lengths are scattered on the lower end.

```
final_tweets <- final_tweets %>% mutate(text_length = nchar(text),
                                          text_word_count = str_count(text, '\\w+'))
```


```
l <- final_tweets %>% filter(experiment_group %in% liberal) %>%
  select(experiment_group, text_length, text_word_count)
```

```
colors <- c("#FDAE61", # Orange
            "#D9EF8B", # Light green
            "#66BD63") # Darker green
```

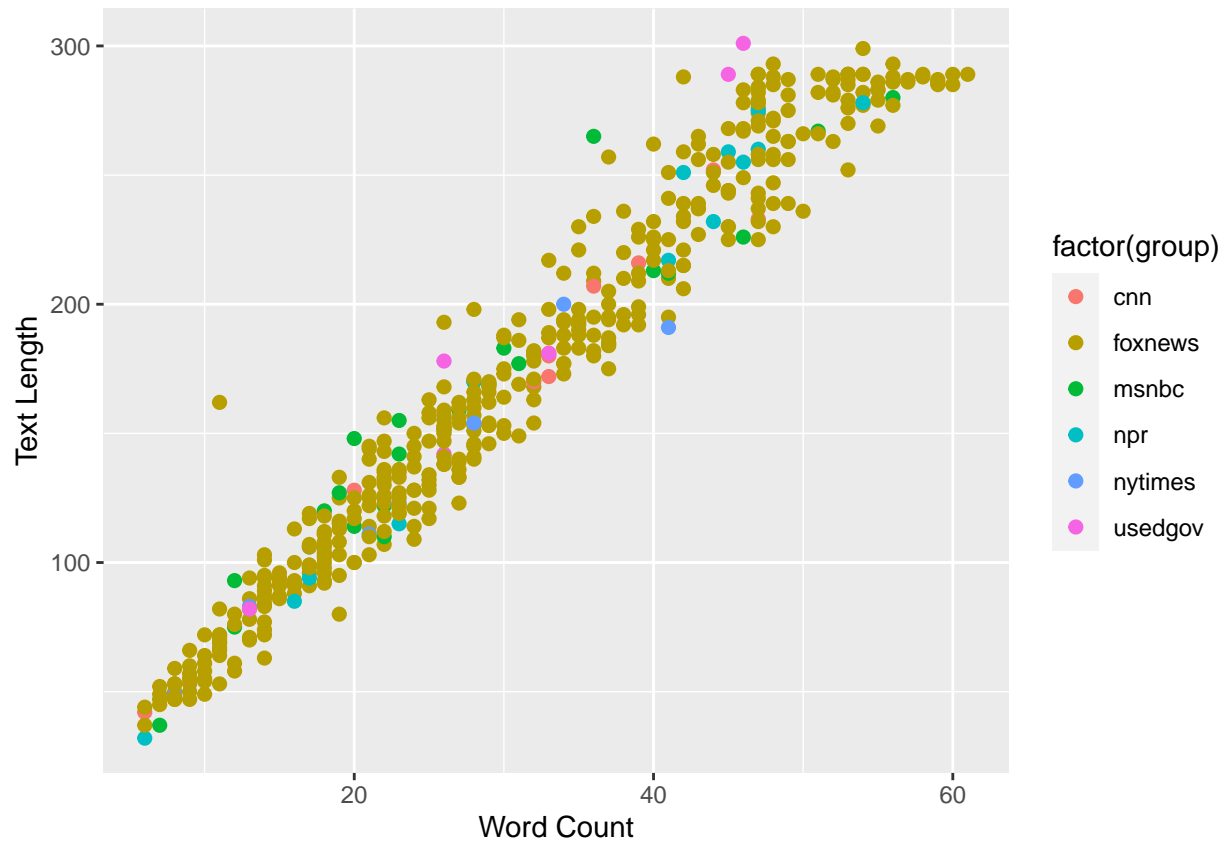
```
x <- l$text_word_count
y <- l$text_length
group <- l$experiment_group
```

Scatter plot

```
ggplot(l, aes(x, y, color = factor(group))) + geom_point(size = 2) + xlab("Word Count") + ylab("Text Length")
```



```
x <- final_tweets$text_word_count
y <- final_tweets$text_length
group <- final_tweets$experiment_group
ggplot(final_tweets, aes(x, y, color = factor(group))) + geom_point(size = 2) + xlab("Word Count") + ylab("Text Length")
```



```
# how does nchar treat emojis? - no.
test = final_tweets[463,] %>% select("text", "text_word_count", "text_length")
```

Tweet Popularity

retweets & experiment_group Which tweet has more *retweets*? Does it happen more often on liberal or conservative outlet? One tweet has 85 retweets, one have 5 retweets, three have 3 retweets but the majority 447 (96%) do not have any retweets.

```
retweet_count <- data.frame(table(final_tweets$retweets)) %>% arrange(desc(Freq))
colnames(retweet_count) <- c("retweets", "freq")
retweet_count
```

```
##   retweets freq
## 1         0 447
## 2         1  16
## 3         2   3
## 4         5   1
## 5        85   1
```

Which outlet has posts with more than 1 retweet? Foxnews and NPR are the two sources where replies have over 1 retweet, with Foxnews holding the highest, 85 retweets.


```
many_retweets <- final_tweets %>% filter(retweets > 1) %>% select(experiment_group, retweets, screen_name)
many_retweets
```

```
##   experiment_group retweets   screen_name
## 1         foxnews      85   RastelliSteve
## 2         foxnews       2     bamakeyman
## 3            npr       5     mkurzawasc
## 4            npr       2     Dollerhide
## 5         foxnews       2   VT_Jeff_RE_Life
```

tweet_likes & experiment_group Which tweet has more likes? Does it happen more often on liberal or conservative outlet? One post has 5446 likes, but the majority (308 out of 478) have 0 likes.

```
tweet_like_count <- data.frame(table(final_tweets$tweet_likes)) %>% arrange(desc(Freq))
colnames(tweet_like_count) <- c("likes_count", "freq")
```

Where is the highest retweet reply? - A tweet addressing FoxNews from someone who is against student loan forgiveness.

```
print(final_tweets[which.max(final_tweets$tweet_likes),]$experiment_group)
```

```
## [1] "foxnews"
```

```
print(final_tweets[which.max(final_tweets$tweet_likes),]$text)
```

```
## [1] "@FoxNews Joe, you cannot spend money without Congress approval. Student loan is not a National "
```

```
print(final_tweets[which.max(final_tweets$tweet_likes),]$tweet_likes)
```

```
## [1] 5446
```

On average, does conservative or liberal sources have more likes and retweets? (after discounting the post with 5446) - NPR has the most average likes and average retweets out of all 5 sources. Replies to Foxnews are 3rd from the bottom in average tweets, even though 87% of the replies in the population belongs to them. On average its replies stand 2nd to last, beating USEdGov, who has less than 1 like on average.

```
no_max_likes <- final_tweets %>% filter(tweet_likes != 5446)
no_max_likes <- no_max_likes %>% group_by(experiment_group) %>%
  summarize(avg_likes = mean(tweet_likes), agg_likes = sum(tweet_likes),
            avg_retweets = mean(retweets), agg_retweets = sum(retweets))
no_max_likes
```

```
## # A tibble: 6 x 5
##   experiment_group avg_likes agg_likes avg_retweets agg_retweets
##   <chr>           <dbl>    <int>         <dbl>         <int>
## 1 cnn             2.25      27         0.0833          1
## 2 foxnews         1.23     502         0.0369         15
## 3 msnbc           1.48      31         0.143           3
## 4 npr             8.23     107         0.615           8
## 5 nytimes         2.5       20          0             0
## 6 usedgov         0.833      5          0             0
```

- When combining the liberal sources, the liberal sources on average have 30% more likes and has 6 times more average retweets than the conservative foxnews.

```
no_max_likes <- final_tweets %>% filter(tweet_likes != 5446) %>%
  mutate(politics = ifelse(experiment_group %in% c('cnn', 'msnbc', 'npr', 'nytimes'), 'liberal',
                           ifelse(experiment_group == 'usedgov', 'controlled', 'conservative')))
no_max_likes <- no_max_likes %>% group_by(politics) %>%
  summarize(avg_likes = mean(tweet_likes), agg_likes = sum(tweet_likes),
            avg_retweets = mean(retweets), agg_retweets = sum(retweets))
no_max_likes
```

```
## # A tibble: 3 x 5
##   politics      avg_likes agg_likes avg_retweets agg_retweets
##   <chr>          <dbl>    <int>      <dbl>      <int>
## 1 conservative    1.23      502      0.0369        15
## 2 controlled      0.833       5       0           0
## 3 liberal         3.43     185      0.222        12
```

User

screen_name

1. Which author has multiple replies? Do they reply to the same source or not?

- 8 people replied twice, 2 of which to multiple news source twitters, but only 1 engage with conservative (FoxNews) and liberal (MSNBC).

```
author_multtweet <- c(data.frame(table(final_tweets$screen_name))) %>% filter(Freq > 1) %>% select(Var1)
author_overlap <- final_tweets %>% filter(screen_name %in% c("DahlmanCarl", "fabulosi_t", "jackSpa81774"))
author_overlap
```

```
##   in_reply_to_screen_name  screen_name statuses_count favourites_count
## 1             MSNBC michael_favreau      30060      16373
## 2             MSNBC RogerWPetersen1      1636       586
## 3             FoxNews thomaslew13      6530        0
## 4             nytimes jackSpa81774793      243        5
## 5             FoxNews thomaslew13      6530        0
## 6             MSNBC michael_favreau      30060      16373
## 7             FoxNews DahlmanCarl      2626      1336
## 8             FoxNews DahlmanCarl      2626      1336
## 9               CNN jackSpa81774793      243        5
## 10            FoxNews PCopposition       59        0
## 11            FoxNews PCopposition       59        0
## 12            usedgov fabulosi_t      5125      5527
## 13            usedgov fabulosi_t      5125      5527
## 14            FoxNews RogerWPetersen1      1636       586
## 15            FoxNews johnbutler410      1637       166
## 16            FoxNews johnbutler410      1637       166
##   followers_count tweet_likes retweets
```

```
## 1      874      0      0
## 2     130      0      0
## 3      12      0      0
## 4       2      0      0
## 5      12      0      0
## 6     874      0      0
## 7       2      1      0
## 8       2      0      0
## 9       2      0      0
## 10      0      1      0
## 11      0      1      0
## 12      72      1      0
## 13      72      0      0
## 14     130      1      0
## 15     184      7      0
## 16     184      0      0
```

`created_at`

Does age of account tell who they might engage with?

```
today <- as.Date("2022-12-08")

ft <- final_tweets %>% mutate(age_dow = substr(created_at, 1, 3)
                             , age_month_day = substr(created_at, 5, 10)
                             , age_time = substr(created_at, 12, 19)
                             , age_yr = substr(created_at, 26, 30),
                             , age_ymd = as.Date(paste0(age_month_day, age_yr), format = "%b %d %Y"),
                             , account_age = today - age_ymd)
```

```
today <- as.Date("2022-12-08")
ft <- ft %>% mutate(account_age = today - age_ymd)

print(paste("min age of acct (days): ", min(ft$account_age)))
```

```
## [1] "min age of acct (days):  5"
```

```
print(paste("max age of acct (days): ", max(ft$account_age)))
```

```
## [1] "max age of acct (days):  5257"
```

```
print(paste("mean age of acct (days): ", mean(ft$account_age)))
```


```
## [1] "mean age of acct (days):  1184.82905982906"
```

```
print(paste("median age of acct (days): ", median(ft$account_age)))
```

```
## [1] "median age of acct (days):  268.5"
```

The youngest `account_age` is 5 days, and the oldest account is 14 years (5257 days)

```
plot(ft$account_age)
```



While most accounts are under 3 years old, there are a handful of accounts in the 4000-5000 days range. Let's look at the text of the accounts with more than 5000 days in age. 6 accounts are over 5000 days old. Majority of them are critical to student loan forgiveness.

One text https://twitter.com/jack_jackson/status/1598689928946323458 @ both NPR and FoxNews. However, the text is an original text (`in_reply_to_status_id` is N/A). Maybe it's okay to keep the `experiment_group` as NPR since mentioning them first prioritize them over Foxnews?

```
ft %>% filter(account_age > 5000) %>% select(experiment_group, text)
```

```
##   experiment_group
## 1      foxnews
## 2      npr
## 3      foxnews
## 4      npr
## 5      foxnews
## 6      foxnews
##
## 1
## 2      @NPR How about, in the meantime, Congress just rewrites the law that makes s
## 3
## 4 @npr @foxnews @dnc @gop Our Constitution requires that all ins and outs of the Treasury originate a
## 5
## 6      @FoxNews Yeah... no... I'm not paying anything extra for that barista
```

description

How many have profile descriptions? More than half of the tweeters don't have an account profile description. Are the share of those with and without description proportional based on who they reply to?

```
(final_tweets %>% mutate(has_profile_desc = ifelse(nchar(description) == 0, 0, 1)) %>% group_by(has_profile_desc))
```

```
## # A tibble: 2 x 2
##   has_profile_desc agg_profile_desc
##           <dbl>         <int>
## 1             0             248
## 2             1             220
```

location

How many have profile location display? Is one location more dense?

Most of the tweets belong to tweeter with no locations (66%).

```
(final_tweets %>% mutate(has_profile_loc= ifelse(nchar(location) == 0, 0, 1)) %>% group_by(has_profile_loc))
```

```
## # A tibble: 2 x 2
##   has_profile_loc agg_profile_loc
##           <dbl>         <int>
## 1             0             312
## 2             1             156
```

ymd & dow


Which day of the week do people discuss student loan forgiveness the most often? Recall that data is from Sunday 11/27 - Tuesday 12/6. Thursdays and Saturdays get the most tweets.

```
final_tweets %>% group_by(dow) %>% summarize(tweet_count = n())
```

```
## # A tibble: 7 x 2
##   dow    tweet_count
##   <chr>         <int>
## 1 Fri             62
## 2 Mon             30
## 3 Sat            114
## 4 Sun             89
## 5 Thu            153
## 6 Tue             14
## 7 Wed             6
```

bar plot of tweet count by day 5 of 9 days have fewer than 20 tweets. Thursday Dec. 1st makes up 33% of all tweets. On Dec 1st, Supreme Court announced they will expedite the process.
- <https://www.nytimes.com/2022/12/01/us/politics/supreme-court-student-loan-forgiveness.html> -
<https://www.washingtonpost.com/politics/2022/12/01/supreme-court-review-student-loan-forgiveness/>

```
ymd_data <- final_tweets %>% group_by(ymd) %>% summarize(tweet_count = n())
barplot(height = ymd_data$tweet_count, names = (ymd_data$ymd))
```



time


What time of day has the most discussion?

```
time <- final_tweets %>% mutate(hour = substr(time, 1, 2))
time_gr <- time %>% group_by(hour) %>% summarize(freq_by_hr = n())
time_gr
```

```
## # A tibble: 24 x 2
##   hour freq_by_hr
##   <chr>      <int>
## 1 00         19
## 2 01         23
## 3 02         16
## 4 03         16
## 5 04         11
## 6 05         13
## 7 06         10
## 8 07          8
## 9 08          2
## 10 09         2
## # ... with 14 more rows
## # i Use 'print(n = ...)' to see more rows
```

bar plot of tweet count by hour Tweets on this topic lulls between 8-10 am. The afternoon has the highest engagement, with a decrease before commuting time, and an rise right after.

```
barplot(height = time_gr$freq_by_hr, names = (time_gr$hour))
```




bar plot with multiple colors for conservative vs. liberal

8pm is a popular time for engagement within our control and liberal groups.

```
final_tweets <- final_tweets %>%
  mutate(hour = substr(time, 1, 2),
         politics = ifelse(experiment_group %in% c('cnn', 'msnbc', 'npr', 'nytimes'), 'liberal',
                          ifelse(experiment_group == 'usedgov', 'controlled', 'conservative')))
stacked_time <- final_tweets %>% group_by(politics, hour) %>% summarize(Freq = n())
```

'summarise()' has grouped output by 'politics'. You can override using the
'.groups' argument.

```
ggplot(stacked_time, aes(fill=politics, y=Freq, x=hour)) +
  geom_bar(position="stack", stat="identity")
```




Is the 8pm mostly due to the Supreme Court announcement on Dec.1st? Yes, Dec 1st makes up over 50% of total tweets at 8pm.

```
dec1_stacked_time <- final_tweets %>% filter(ymd == "2022-12-01") %>%
  group_by(politics, hour) %>% summarize(Freq = n())
```

```
## 'summarise()' has grouped output by 'politics'. You can override using the
## '.groups' argument.
```

```
ggplot(dec1_stacked_time, aes(fill=politics, y=Freq, x=hour)) +
  geom_bar(position="stack", stat="identity")
```

User Popularity

favourites_count Which media sources has engagement from the most favoured tweeter?

```
fav_counts_tweet <- data.frame(table(final_tweets$favourites_count)) %>% arrange(desc(Freq))
```

followers_count Which media sources has engagement from the tweeters with the highest & lowest follower count??

verified Are there any verified accounts? If so, where did they engage with? None of the author is verified.

```
unique(final_tweets$verified)
```

```
## [1] "False"
```

Is there an reason why someone who is against student loan forgiveness would reply to fox vs. cnn or vice versa? - refer to max like, why is the opposition not addressing someone?

EDA - with annotations

Give summary of the stand

What is the stand of the “older twitter accounts”?

*Stacked bar plot on views on student loans and how it matches up against **experiment_group*** `gg-`
`plot(stacked_time, aes(fill=politics, y=Freq, x=hour)) + geom_bar(position="stack", stat="identity")`